



POC Environment Prerequisites

Client Preparation Checklist

What to prepare before AXIA deploys the proof-of-concept

axiasecurity.io • hello@axiasecurity.io

What the client team prepares before Axia deploys the proof-of-concept.

Axia ships a two-tier engine — a local Gatekeeper plus an LLM judgment endpoint — into the client's own environment. Content never leaves the client perimeter. This checklist covers what the **client** provisions; Axia provisions the platform layer on top of it via Terraform and completes the deployment remotely.

***Reference deployment:** AWS + Amazon Bedrock. The same structure applies to the Azure (Azure OpenAI Service) and GCP (Vertex AI) variants using the equivalent managed services. Axia supplies the cloud-specific values for your environment during onboarding.*

1. Cloud Infrastructure

***Note:** All infrastructure in this section is provisioned by Axia via Terraform. The client only needs to provide the cloud account access and network (VPC / subnet) details.*

Item	Spec (AWS reference)
Compute node 1 — Platform	m5.2xlarge minimum (m5.3xlarge recommended) · 100 GB SSD · Private subnet
Compute node 2 — Inference/Compute	c7i-flex.4xlarge · 200 GB SSD · Same private subnet
Document database	DocumentDB db.r6g.large or larger · TLS enabled · port 27017 reachable from compute nodes
Object storage	Private bucket, encrypted at rest (SSE-S3 or equivalent), with scoped credentials
Log group	Centralized log group, e.g. <code>/platform/services</code>

Both compute nodes must sit in **private subnets with no public IP**.

2. Host Software (on both compute nodes)

- Docker Engine **v28.0+**
- Docker Compose **v2.40+**
- Cloud CLI **v2** (AWS CLI v2 for the reference deployment)

3. Cloud IAM Permissions

The compute-node instance role must grant:

- **Model inference** — invoke + streaming-invoke on the managed LLM service (`bedrock:InvokeModel`, `bedrock:InvokeModelWithResponseStream` for the AWS reference; the Azure / GCP equivalents for those variants)
- **Logging** — create log streams and put log events (`logs:CreateLogStream`, `logs:PutLogEvents`, `logs:DescribeLogGroups`, `logs:DescribeLogStreams`)
- **Instance management** — the cloud's managed instance-access policy (e.g. `AmazonSSMManagedInstanceCore`) so Axia can perform remote deployment
- **Secrets / parameters** — read the platform's startup secrets and config from the parameter store, scoped to the platform path (e.g. `ssm:GetParameter`, `ssm:GetParameters` on `/platform/*`)
- **Container image pull** — cross-account pull from Axia's container registry (`ecr:GetAuthorizationToken`, `ecr:BatchGetImage`, `ecr:GetDownloadUrlForLayer`, `ecr:BatchCheckLayerAvailability`). **Axia provides the registry account ID and region at onboarding.**

Object-storage access: a dedicated service identity with put / get / delete / list scoped to the platform bucket.

Model access must be enabled in the cloud console before deployment (for AWS Bedrock + Anthropic Claude, complete the First-Time-Use form in the AWS console).

4. Network / Firewall

Rule	Detail
Outbound 443	Managed cloud services (container registry, model inference, object storage, logs) + <code>googleapis.com</code> (for the Google Workspace integration)
Inbound — Dashboard	Dashboard port (reference: 5173) — from corporate network / VPN only
Inbound — Database	Document database port (27017) — from the compute-node security group only
Deployment access	Cloud session-manager (preferred) or SSH port 22 restricted to Axia's source IP

No service ports should be exposed to the public internet.

5. Google Workspace / GCP Integration (*required for email + directory ingestion*)

Complete the following in GCP and the Google Workspace Admin console:

1. **GCP project** with these APIs enabled: Gmail, Admin SDK, Google Drive, IAM Service Account Credentials
2. **Workload Identity Federation (WIF)** pool + provider linked to the client's cloud account
3. **Service account** (e.g. `axia-drive`) with Domain-Wide Delegation enabled
4. **WIF grant**: Service Account Token Creator role for the compute-node instance-role principal
5. **Workspace Admin authorization**: grant the service account read scopes for Gmail (read), Admin Directory (read), Drive (read)
6. **Impersonation user**: a Google Workspace mailbox for the platform to act as (no special privileges required)

Deliverables back to Axia: GCP project ID · service-account email · WIF provider resource name · impersonation user email

6. Optional Integrations

Integration	What's needed
DLP signal source (e.g. Netskope)	Tenant URL + read-only API token (scoped ingestion role)
Send-to-reviewer email	Mailbox app password (16-char) + reviewer's email address

7. What Axia Provides

- Container-registry credentials + the registry account/region for pulling platform software
- The exact WIF principal string for the Google service-account grant
- Terraform-based provisioning of the platform layer
- Remote deployment — approximately 2–3 hours once all prerequisites above are met

8. Readiness Checklist

- ☐ Two private-subnet compute nodes provisioned to spec, no public IP
- ☐ Document database reachable from the compute nodes over TLS
- ☐ Private, encrypted object-storage bucket with scoped credentials
- ☐ Instance role carries inference, logging, instance-management, secrets, and registry-pull permissions
- ☐ Managed LLM model access enabled in the cloud console
- ☐ Firewall rules applied (outbound 443; inbound dashboard + database restricted; deployment access path open to Axia)
- ☐ Google Workspace / GCP integration complete, four deliverables sent to Axia
- ☐ Optional integrations decided (DLP source, send-to-reviewer email)

End of Document

AXIA Security • axiasecurity.io • hello@axiasecurity.io